

CS60002: Distributed Systems
Spring 2025
Assignment – 1

Due: February 5, 2025 (submit in moodle)

Answers must be typed mostly and finally a single pdf document should be submitted. You can insert scans of hand-drawn pictures if really needed.

1. Write clear and structured pseudocode for all algorithms, in the form discussed in class.
 2. Use concise yet descriptive variable and function names.
 3. Avoid implementation-specific syntax (e.g., Python, C++).
 4. Include comments in the pseudocode to explain key steps.
 5. Beside the variables mentioned in a problem (if any), you can add any other variable you want.
 6. If you make any other assumptions (over and above what is given) in a problem, please write them clearly at the beginning of your answer.
-
1. Consider a system with a clock that runs slow with a drift rate of 10^{-5} . Cristian's algorithm is used to synchronize the clock from a time server every 30 minutes. To estimate the delay in getting the time, a simple algorithm is used – the system just measures the round trip delay and divides it by 2. The maximum delay of the link between the system and the time server is 2 milliseconds (one way). The processing time at the client after receiving the time from the server is 1 millisecond. Assume that the time server's clock has zero drift. **What is the maximum possible difference between the two clocks (the time server and the client system) just before and just after a synchronization?** Show all calculations.
 2. Suppose that 3 nodes always have their clocks synchronized to within 1 second of each other. A simple global snapshot algorithm is proposed in which each node takes its local snapshot at the hour, i.e., at 9 AM, 10 AM, ... as per its local clock; no other messages are exchanged. If we take a global state to be the set of local snapshots of the 3 nodes taken at say 10 AM, is it guaranteed to be a consistent global state? Justify your answer clearly (Prove if yes, give a counterexample if no).
 3. Write a distributed algorithm for constructing a spanning tree rooted at a specific node in an asynchronous, reliable system. The root node will initiate the algorithm. At the termination of the algorithm, each node should know its parent and children set in the spanning tree. In addition, the root node should know that the spanning tree has been constructed (i.e., the algorithm has terminated). The model is asynchronous, reliable communication, arbitrary topology. **(For your practice, already discussed in class, this will not be graded).**
 4. Suppose you are given a connected, undirected graph with a special node i and an upper bound T on the message delay over any link. Design a global state collection algorithm that does not rely on FIFO channels and does not use piggybacking. At the end of your algorithm, the global state should be available at node i and i should know that state collection is complete. **Analyze its message complexity.** Clocks in the nodes are not synchronized, but you can assume that they have no drift. Assume that every process knows an upper bound on the no. of processes in the system, and processing times at nodes are negligible. Communication is reliable.

5. Design a permission based distributed mutual exclusion algorithm in which a node will require 0 messages per critical section entry in the best case. However, in the worst case, it will require $2(n-1)$ message per critical section entry similar to Ricart-Agarwala's algorithm. **Identify what the best and the worst cases are.** The proposed algorithm need not maintain the fairness property of Ricart-Agarwala's algorithm.
6. Design an asynchronous distributed algorithm to find a BFS spanning tree from a designated node X to all other nodes in an undirected, connected graph. X will be the only node to initiate the algorithm, and no flooding should be used. At the end of the algorithm, each node should have a local variable *Level* whose value set to its shortest path length from X . The BFS spanning tree should be built level by level, i.e., nodes at any level x should be added to the tree only after all nodes at level $x-1$ have been added to the tree (X should synchronize this). **What is the message complexity of the algorithm?** Do not worry about how X gets to know when all the values are correctly computed at the nodes (i.e., about detecting termination of the algorithm at X). Communication is reliable.
7. Design an asynchronous distributed algorithm to find a DFS spanning tree in an undirected, connected graph. Nodes have unique id. The algorithm may be started by more than one node acting as the initiator (root), each starting independently at any time. However, if the spanning tree building of one or more initiators overlap in time, finally only one DFS spanning tree should exist, specifically the one built by the initiator with the maximum id among all such initiators.
8. Design a token based mutual exclusion algorithm that uses a rooted spanning tree to pass tokens along the edges of the spanning tree. The spanning tree is defined by a "*to_token*" variable at each node which points to its parent in the tree, the value for the initial root is set to itself. The token is initially at the root of the spanning tree built. The root is allowed to be changed over time if needed (while maintaining a proper rooted tree), but the underlying unrooted tree remains the same. You do not have to construct the spanning tree, assume that it exists (i.e., the *to_token* variables are set properly). If there is no pending request in the system, the token should stay with the node that used it last. A node that wants to use the token should have to request for it when it needs it. Each node in the tree should be in charge of requesting for and use of the token on behalf of all nodes in the subtree rooted at it, including itself. Use the ideas discussed in class to write the algorithm clearly. **What do you think is the no. of messages needed for CS entry for your algorithm?** Model is asynchronous, reliable, nodes have unique ids.