

Distributed Computing Concepts - Global State in Distributed Systems

Prof. Nalini Venkatasubramanian
230 Distributed Systems - Week 3

**-includes slides/examples from
Indy Gupta (UIUC), Coulouris(book) and
Kshemkalyani&Singhal (book slides)**

Why Global State?

- Distributed applications/services execute concurrently on multiple machines.
- A **Snapshot** of the distributed application, i.e. a **global picture is useful**
 - *Checkpointing*: can restart distributed application on failure
 - *Garbage collection* of objects: objects at servers that don't have any other objects (at any servers) with pointers to them
 - Deadlock detection: Useful in database transaction systems
 - Termination of computation: Useful in batch computing systems like Folding@Home, SETI@Home



The Power of Snapshots
Stateful Stream Processing
with Apache Flink

Stephan Ewen
dataArtisans

QCon San Francisco, 2017

What constitutes global state?

- **Global Snapshot = Global State**

Individual state of *each process* in the distributed system

+


Individual state of *each communication channel* in the distributed system

- Capture the **instantaneous** state of each process
- Capture the instantaneous state of each communication channel, i.e., *messages* in transit on the channels


Simulate A Global State

- The notions of global time and global state are closely related.
 - But, merely synchronizing clocks and taking local snapshots is not enough
 - Need to account for messages in transit
- A process can (without *freezing* the whole computation) compute the *best possible approximation* of a global state [Chandy & Lamport 85]
- A global state that *could* have occurred
 - No process in the system can decide whether the state did really occur
 - Guarantee stable properties (*i.e.* once they become true, they remain true)


Event Diagram




Equivalent Event Diagram



Poset Diagram




Rubber Band Transformation



Consistent Cuts

- A cut (or time slice) is a zigzag line cutting a time diagram into 2 parts (past and future)
 - E is augmented with a cut event c_i for each process $P_i: E' = E \cup \{c_i, \dots, c_n\} \therefore$
 - A cut C of an event set E is a finite subset $C \subseteq E: e \in C \wedge e' <_i e \rightarrow e' \in C$
 - A cut C_1 is later than C_2 if $C_1 \supseteq C_2$
 - A consistent cut C of an event set E is a finite subset $C \subseteq E : e \in C \wedge e' < e \rightarrow e' \in C$
 - i.e. a cut is consistent if every message received was previously sent (but not necessarily vice versa!)

Cuts (Summary)



Consistent Cuts

- Some Theorems
 - For a consistent cut consisting of cut events c_1, \dots, c_n , no pair of cut events is causally related. i.e. $\forall c_i, c_j \sim (c_i < c_j) \wedge \sim (c_j < c_i)$
 - For any time diagram with a consistent cut consisting of cut events c_1, \dots, c_n , there is an equivalent time diagram where c_1, \dots, c_n occur simultaneously. i.e. where the cut line forms a straight vertical line
 - All cut events of a consistent cut *can occur simultaneously*

Global States of Consistent Cuts

- The global state of a distributed system is a collection of the local states of the processes and the channels.
- A *global state* computed along a consistent cut is **correct**
- The *global state* of a consistent cut comprises the local state of each process at the time the cut event happens and the set of all messages sent but not yet received
- The *snapshot problem* consists in designing an efficient protocol which yields only consistent cuts and to collect the local state information
 - Messages crossing the cut must be captured
 - Chandy & Lamport presented an algorithm assuming that message transmission is FIFO

Distributed Global Snapshot: Requirements

- Snapshot should not interfere with normal application actions, and it should not require application to stop sending messages
- Each process is able to record its own state
 - Process state: Application-defined state or, in the worst case:
 - its heap, registers, program counter, code, etc. (essentially the coredump)
- Global state is collected in a distributed manner
- Any process may initiate the snapshot
 - Assume just one snapshot run for now

System Model for Global Snapshots

- The system consists of a collection of n processes p_1, p_2, \dots, p_n that are connected by channels.
- There are no globally shared memory and physical global clock and processes communicate by passing messages through communication channels.
- C_{ij} denotes the channel from process p_i to process p_j and its state is denoted by SC_{ij} .
- The actions performed by a process are modeled as three types of events:
 - Internal events, the message send event and the message receive event.
 - For a message m_{ij} that is sent by process p_i to process p_j , let $\text{send}(m_{ij})$ and $\text{rec}(m_{ij})$ denote its send and receive events.

Process States and Messages in transit

- At any instant, the state of process p_i , denoted by LS_i , is a result of the sequence of all the events executed by p_i till that instant.
 - For an event e and a process state LS_i , $e \in LS_i$ iff e belongs to the sequence of events that have taken process p_i to state LS_i .
 - For an event e and a process state LS_i , $e \notin LS_i$ iff e does not belong to the sequence of events that have taken process p_i to state LS_i .
- For a channel C_{ij} , the following set of messages can be defined based on the local states of the processes p_i and p_j
Transit: $\text{transit}(LS_i, LS_j) = \{m_{ij} \mid \text{send}(m_{ij}) \in LS_i \vee \text{rec}(m_{ij}) \notin LS_j\}$

Chandy-Lamport Distributed Snapshot Algorithm

- **Assumes FIFO communication in channels**
- Uses a control message, called a *marker* to separate messages in the channels.
 - After a site has recorded its snapshot, it sends a marker, along all of its outgoing channels before sending out any more messages.
 - The marker separates the messages in the channel into those to be included in the snapshot from those not to be recorded in the snapshot.
- **When to capture local state?**
 - A process must record its snapshot no later than when it receives a marker on any of its incoming channels.
- **Termination**
 - The algorithm terminates after each process has received a marker on all of its incoming channels.
- All the local snapshots get disseminated to all other processes and all the processes can determine the global state.

Chandy-Lamport Distributed Snapshot Algorithm

Marker receiving rule for Process P_i

If (P_i has not yet recorded its state) *it*

records its process state now

records the state of c as the empty set

turns on recording of messages arriving over other channels

else

P_i records the state of c as the set of messages received over c
since it saved its state

Marker sending rule for Process P_i

After P_i has recorded its state, for each outgoing channel c :


P_i sends one marker message over c

(before it sends any other message over c)

Snapshot Example

e_1^3

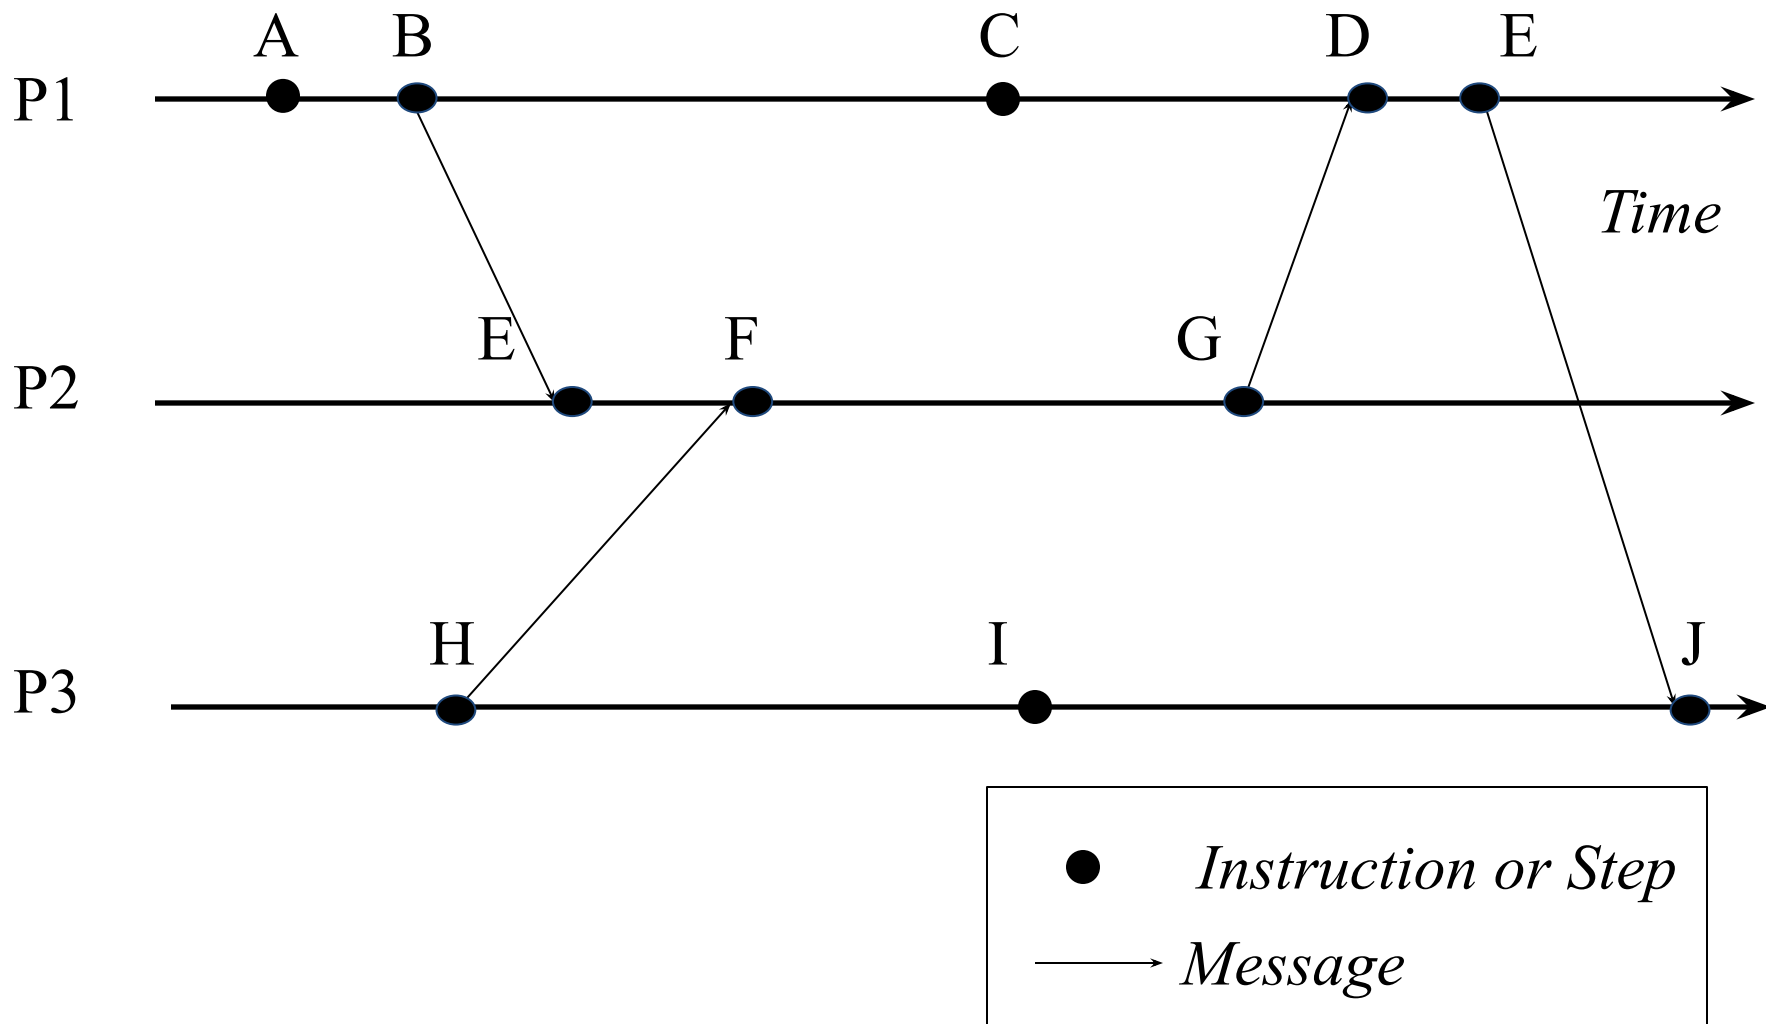
From: Indranil Gupta



1. P1 initiates snapshot: records its state (S_1); sends Markers to P2 & P3; turns on recording for channels C21 and C31
- 2- P2 receives Marker over C12, records its state (S_2), sets $\text{state}(C12) = \{\}$ sends Marker to P1 & P3; turns on recording for channel C32
- 3- P1 receives Marker over C21, sets $\text{state}(C21) = \{a\}$
- 5- P2 receives Marker over C32, sets $\text{state}(C32) = \{b\}$
- 6- P3 receives Marker over C23, sets $\text{state}(C23) = \{\}$
- 7- P1 receives Marker over C31, sets $\text{state}(C31) = \{\}$


Snapshot Example


From: Indranil Gupta (CS425 - Distributed Systems course, UIUC)




P1 is Initiator:


- Record local state S1,
- Send out markers
- Turn on recording on channels C_{2I} , C_{3I}








- First Marker!
- Record own state as S3
- Mark C_{13} state as empty
- Turn on recording on other incoming C_{23}
- Send out Markers














- Duplicate!
- $C_{2l} = \langle \text{message } G \square D \rangle$







Algorithm has Terminated




Collect the Global Snapshot Pieces



Our Global Snapshot Example ...



... is causally correct



Chandy-Lamport Extensions: Spezialetti-Kerns and others

- Exploit concurrently initiated snapshots to reduce overhead of local snapshot exchange
- Snapshot Recording
 - Markers carry identifier of initiator – first initiator recorded in a per process "*master*" variable.
 - **Region** - all the processes whose master field has same initiator.
 - Identifiers of concurrent initiators recorded in "*id-border-set*."
- Snapshot Dissemination
 - Forest of spanning trees is implicitly created in the system. Every Initiator is root of a spanning tree; nodes relay snapshots of rooted subtree to parent in spanning tree
 - Each initiator assembles snapshot for processes in its region and exchanges with initiators in adjacent regions.
- Others: multiple repeated snapshots; wave algorithm

Computing Global States without FIFO Assumption

- In a non-FIFO system, a marker cannot be used to delineate messages into those to be recorded in the global state from those not to be recorded in the global state.
- In a non-FIFO system, either some degree of inhibition or piggybacking of control information on computation messages to capture out-of-sequence messages is required

Computing Global States without FIFO Assumption - Lai-Yang Algorithm

- Uses a *coloring* scheme that works as follows
 - White (before snapshot); Red (after snapshot)
 - Every process is initially white and turns *red* while taking a snapshot. The equivalent of the “Marker Sending Rule” (virtual broadcast) is executed when a process turns *red*.
 - Every message sent by a white (*red*) process is colored white (*red*).
 - Thus, a white (*red*) message is a message that was sent before (after) the sender of that message recorded its local snapshot.
 - Every white process takes its snapshot at its convenience, but no later than the instant it receives a *red* message.

Computing Global States without FIFO Assumption - Lai-Yang Algorithm (cont.)

- Every white process records a history of all white messages sent or received by it along each channel.
- When a process turns red, it sends these histories along with its snapshot to the initiator process that collects the global snapshot.
- Determining Messages in transit (i.e. White messages received by red process)
 - The initiator process evaluates $\text{transit}(\text{LS}_i, \text{LS}_j)$ to compute the state of a channel C_{ij} as given below:
 - $\text{SC}_{ij} = \{\text{white messages sent by } p_i \text{ on } C_{ij} - \text{white messages received by } p_j \text{ on } C_{ij}\}$
 - $= \{ \text{send}(M_{ij}) | \text{send}(m_{ij}) \in \text{LS}_i \} - \{ \text{rec}(m_{ij}) | \text{rec}(m_{ij}) \in \text{LS}_j \}.$

Computing Global States without FIFO Assumption: Termination

● First method

- Each process P_i keeps a counter $cntri$ that indicates the difference between the number of white messages it has sent and received before recording its snapshot, i.e number of messages still in transit.
- It reports this value to the initiator along with its snapshot and forwards all white messages, it receives henceforth, to the initiator.
- Snapshot collection terminates when the initiator has received $\sum_i cntri$ number of forwarded white messages.

● Second method

- Each red message sent by a process piggybacks the value of the number of white messages sent on that channel before the local state recording. Each process keeps a counter for the number of white messages received on each channel.
- Termination – Process P_i receives as many white messages on each channel as the value piggybacked on red messages received on that channel.

Computing Global States without FIFO Assumption: Mattern's Algorithm

- Uses Vector Clocks
 - All process agree on some future **virtual time** s or a **set of virtual time instants** s_1, \dots, s_n which are mutually concurrent and **did not yet** occur
 - A process takes its local snapshot at **virtual time** s
 - After time s the local snapshots are collected to construct a global snapshot
 - P_i ticks and then fixes its next time $s = C_i + (0, \dots, 0, 1, 0, \dots, 0)$ to be the common snapshot time
 - P_i broadcasts s
 - P_i blocks waiting for all the acknowledgements
 - P_i ticks again (setting $C_i = s$), takes its snapshot and broadcast a dummy message (i.e. force everybody else to advance their clocks to a value $\geq s$)
 - Each process takes its snapshot and sends it to P_i when its local clock becomes $\geq s$

Computing Global States without FIFO Assumption (Mattern cont)

- Inventing a $n+1$ **virtual** process whose clock is managed by P_i
- P_i can use its clock and because the virtual clock C_{n+1} ticks only when P_i initiates a new run of snapshot :
 - The first n components of the vector can be omitted
 - The first broadcast phase is unnecessary
 - Counter modulo 2
- Termination
 - Distributed termination detection algorithm [**Mattern 87**]

Optional video: has detailed example illustrating the challenge of capturing global snapshots.

<https://www.youtube.com/watch?v=ao58xine3jM>