


# Handling Deadlock




# Handling Deadlock

Issues:

- Reusable vs. consumable resources
- Resource vs. communication deadlock
- AND vs OR deadlock
- Wait-For graphs (WFG)
- Prevention, Avoidance, Detection??
- Resolution??

# Deadlock Detection Strategies



## Requirements:


- No undetected deadlocks
- No “false” or “phantom” deadlocks
  - Detecting a deadlock even when there is none present in the system

## Strategies:

- Centralized
- Distributed
  - Path-pushing vs. Edge-chasing algorithms
- Hierarchical

# A Simple Centralized Algorithm

## (Ho-Ramamoorthy)




- Each node has a status table, contains status (resources locked and resources waited on) of all processes at that node
- A central site periodically collects the status table from all nodes, constructs the WFG and checks for cycles
- If no cycle detected, no deadlock
- If cycle detected, status from all nodes requested again and WFG constructed using ONLY information common both times. If the same cycle is detected again, deadlock is declared.
- Does NOT work!! Why??
- Can you make it work with additional information?


# Chandy-Misra-Haas Algorithm for AND Deadlocks



- Distributed control
- An “Edge-Chasing” algorithm
- Uses a special probe message of the form  $(i, j, k)$  where:
  - $p_i$  : process originally initiating deadlock detection
  - $p_j$  : current sender
  - $p_k$  : destination/receiver
- A process  $p_i$  is dependant on another process  $p_j$  if there exists a path from  $p_i$  to  $p_j$  in the WFG
- If  $p_i$  and  $p_j$  are in the same node,  $p_i$  is locally dependent on  $p_j$



- Main Idea:
  - A blocked process  $p_i$  initiates detection by sending probes to all processes  $p_k$  at another node on which it is dependent (directly or indirectly)
  - A process receiving a probe  $(i, j, k)$  forwards it to all processes it is waiting for after changing the  $j$  and  $k$  fields appropriately,  $i$  remains unchanged.
  - Thus the probe message travels across the edges of the WFG; if it comes back to the initiator, WFG has a cycle and we have a deadlock.
- Each process maintains an array *dependent<sub>j</sub>*:  
*dependent<sub>j</sub>(j)* is true if  $P_i$  knows that  $P_j$  is dependent on it. (initially set to false for all  $i$  &  $j$ ).



# The Algorithm

Sending the probe (from  $P_i$ ) :

if  $P_i$  is locally dependent on itself then deadlock.

else for all  $P_j$  and  $P_k$  such that

(a)  $P_i$  is locally dependent upon  $P_j$ , and

(b)  $P_j$  is waiting on  $P_k$ , and

(c)  $P_j$  and  $P_k$  are on different sites, send probe( $i, j, k$ )  
to the home site of  $P_k$ .

Receiving the probe ( $i, j, k$ ) at  $P_k$  :

if (d)  $P_k$  is blocked, and

(e)  $dependent_k(i)$  is false, and


(f)  $P_k$  has not replied to all requests of  $P_j$ ,

then begin

$dependent_k(i) := \text{true};$

if  $k = i$  then  $P_i$  is deadlocked

else ...



Receiving the probe (contd.):

.....

else for all  $P_m$  and  $P_n$  such that

(a')  $P_k$  is locally dependent upon  $P_m$ , and

(b')  $P_m$  is waiting on  $P_n$ , and

(c')  $P_m$  and  $P_n$  are on different sites,

send probe( $i, m, n$ ) to the home site of  $P_n$ .

end.



# Example

