

Memcached: A Distributed Memory Object caching System

History

- Built by Brad Fitzpatrick;s company, Danga Interactive, around 2003, for LiveJournal, a community based journaling platform
- First described in a Linux Journal paper “Distributed caching with Memcached” in 2004
- Used extensively by almost all big companies like Facebook. Twitter (X), Amazon. Google,. Microsft....in some form or other
- Available for download as open source software from <https://memcached.org>

What is memcached

- High performance, distributed memory object caching system
- Designed to reduce load on backend DBs
 - In memory caching of dynamic DB contents
- Pools together spare memory from your servers (or more often, has dedicated memcached servers for caching only) to create a distributed cache
- What does it cache?
 - (key, value) pairs, along with an expiration time and some other optional flags
 - Value is a raw binary (serialized) data, does not understand any complex types/structures
- Main goal
 - $O(1)$ access to data
 - Read latency of < 1 ms
 - High end memcached servers able to serve millions of queries per second

Broad Architecture

- A set of memcached servers acting as cache
- A set of clients who wish to store/retrieve data from the cache
- A memcached implementation consists of
 - A client part to choose which server to read and write for a specific data item
 - A server part to store and retrieve items, and decide when to evict items from cache to reuse
- Servers are disconnected from each other
 - Completely unaware of each other, no synchronization needed, no replication (in basic form)
 - Simple cache invalidation
 - Data deleted/overwritten in the server holding it directly by client
- Simple APIs to get/store/replace/delete data items

Eviction/reuse

- Cache is allowed to “forget”
- Data item evicted if
 - Expiration time over
 - Cache is full and new item needs to be added
- LRU based eviction policy

Accessing Data

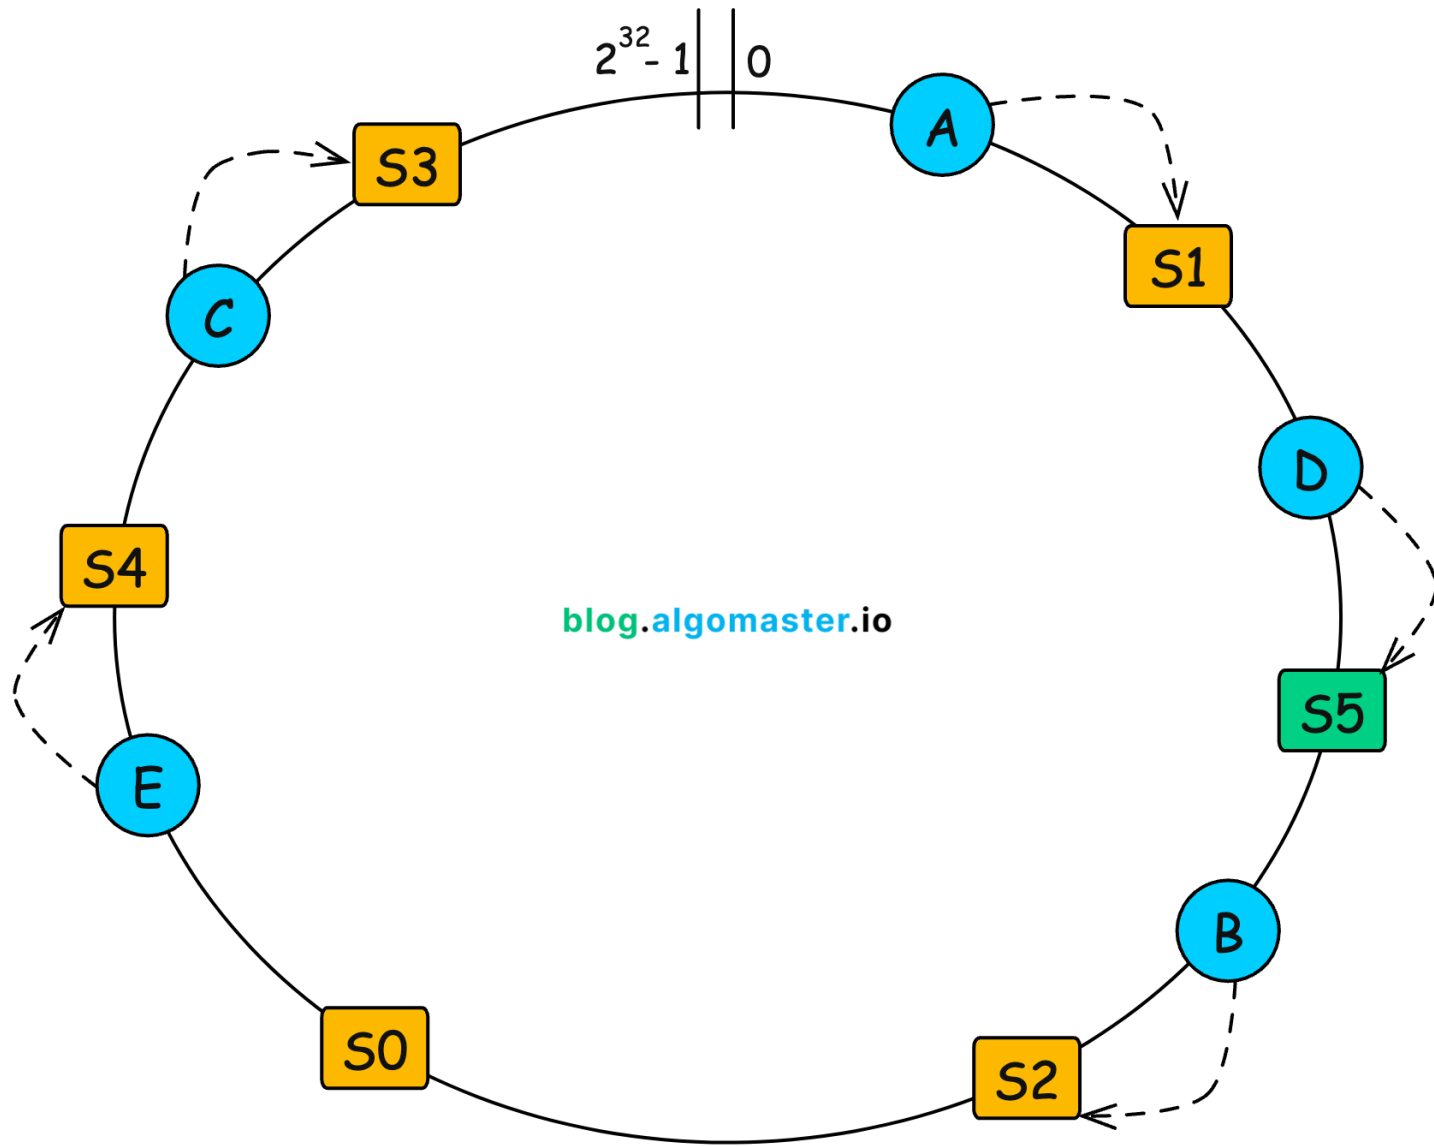
- Clients know the list of servers involved in storing data of a service (IP, port)
- Mapping data to server
 - Use hashing, just two-step
 - One to find the server
 - One to find data item within server
 - Hash the key in the (key, value) pair using
 - $H(\text{key}) \% \text{no_of_servers}$
 - Connect to the server to get the data
- Two completely different keys can map to the same server
 - Server keeps the table of (key, value) pairs mapped to it in an internal hash table
 - 2nd level of hashing

Problem

- What happens when a server joins or fails?
 - Hashing may need to be done again for all items, which is very expensive
- One scenario
 - Fails but a spare server put up
 - Same also for taking a server down temporarily for update
 - Memcached provides mechanism for backup/restore of the cache
- Still does not solve cases when a server stays down or a new server joins
 - No. of servers change


Consistent Hashing

- Hash both the keys and the servers onto a given range
- Consider the range as a circular ring (Hash ring) with smallest and largest value tied at the end
- Each key and server will be a point on this ring
- To find the server storing a key, start from the point corresponding to the key, go clockwise (or anticlockwise, same for all keys) on the ring, and find the first server hit
- Store the key in that server
- Advantage
 - Deleting a server only affects the keys hashed in the range stored in that server
 - Similar for adding a new server
 - No need to rehash any other key in any other server



What is it good for?

- Modern DC networks are very fast (a few microsecond latency)
- However, data read delays from DRAM are still much faster
- Means that memcached is not good for caching small objects, like file blocks
 - Network delay of getting to the server will be significant compared to data transfer time
 - The delay will start to show!
 - Also, file systems typically have their own cache
- Good for large objects, like images, videos, large webpages
 - Overhead of accessing server is small compared to the actual data transfer time



- An example use case
 - An application needs to resize photos for different screen sizes
 - Option 1: resize in specific device
 - Energy-intensive for smaller devices
 - Option 2: access from source, resizing each time
 - Large resizing latency
 - Use memcached
 - Resize in cloud, store in memcached and access

- Note that memcached is
 - A caching system, not a storage system, so not persistent
 - Evicts data items, so does not guarantee that a get later gets what is stored earlier even within expiration period
 - In its simplest form, not replicated (to avoid inter-server synchronization overhead) so data items may be lost on failure
 - Ok as it is just a cache
 - Can do replication also

Locating servers

- Assumed that all clients know all memcached server IPs for a specific service
- Does not scale. So?